

Using the dplyr package

Department of Biostatistics and Bioinformatics

Steve Pittard wsp@emory.edu

February 1, 2018

Data Frames



Why Use Data Frames ?

- A data frame is a special type of list that contains data in a format that allows for easier manipulation, reshaping, and open-ended analysis
- Data frames are tightly coupled collections of variables. It is one of the more important constructs you will encounter when using R so learn all you can about it
- A data frame is an analogue to the Excel spreadsheet but is much more flexible for storing, manipulating, and analyzing data
- Data frames can be constructed from existing vectors, lists, or matrices. Many times they are created by reading in comma delimited files, (CSV files), using the `read.table` command
- Once you become accustomed to working with data frames, R becomes so much easier to use

Why Use Data Frames ?

Use the **dataframe()** function to create a data frame. It looks like a matrix but allows for mixed data types

```
names <- c("P1","P2","P3","P4","P5")
temp  <- c(98.2,101.3,97.2,100.2,98.5)
pulse <- c(66,72,83,85,90)
gender <- c("M","F","M","M","F")
```

```
my_df <- data.frame(names,temp,pulse,gender) # Much more flexible
```


	names	temp	pulse	gender
1	P1	98.2	66	M
2	P2	101.3	72	F
3	P3	97.2	83	M
4	P4	100.2	85	M
5	P5	98.5	90	F

```
plot(my_df$pulse ~ my_df$temp,main="Pulse Rate",xlab="Patient",ylab="BPM")
mean(my_df[,2:3])
temp pulse
99.08 79.20
```

Why Use Data Frames ?

Once you have a data frame you could edit it with the Workspace viewer in RStudio although this doesn't generalize. Imagine if your data set had 10,000 lines ?

```
data(mtcars) # Load the builtin mtcars dataframe
```

The screenshot shows the RStudio Environment pane. At the top, there are tabs for 'BIOS60R.Rstudio.R', 'chicagocrime.R', 'twitter.R', and 'mtcars'. The 'mtcars' tab is selected. Below the tabs, there is a table with 32 observations of 11 variables. The table has columns for row.names, mpg, cyl, disp, hp, drat, wt, qsec, vs, am, gear, and carb. The data is as follows:

row.names	mpg	cyl	disp	hp	drat	wt	qsec	vs	am	gear	carb
1 Mazda RX4	21.0	6	160.0	110	3.90	2.620	16.46	0	1	4	4
2 Mazda RX4 Wag	21.0	6	160.0	110	3.90	2.875	17.02	0	1	4	4
3 Datsun 710	22.8	4	108.0	93	3.85	2.320	18.61	1	1	4	1
4 Hornet 4 Drive	21.4	6	258.0	110	3.68	3.215	19.44	1	0	3	1
5 Hornet Sportabout	18.7	8	360.0	175	3.15	3.440	17.82	0	0	3	2
6 Valiant	18.1	6	225.0	105	2.76	3.460	20.22	1	0	3	1
7 Duster 360	14.3	8	360.0	245	3.21	3.570	15.84	0	0	3	4
8 Merc 240D	24.4	4	146.7	62	3.69	3.190	20.00	1	0	4	2
9 Merc 230	22.8	4	140.8	95	3.92	3.150	22.90	1	0	4	2
10 Merc 280	19.2	6	167.6	123	3.92	3.440	18.30	1	0	4	4
11 Merc 280C	17.8	6	167.6	123	3.92	3.440	18.90	1	0	4	4
12 Merc 450SE	16.4	8	275.8	180	3.07	4.070	17.40	0	0	3	3
13 Merc 450SL	17.3	8	275.8	180	3.07	3.730	17.60	0	0	3	3
14 Merc 450SLC	15.2	8	275.8	180	3.07	3.780	18.90	0	0	3	3
15 Cadillac Fleetwood	16.4	8	472.0	205	2.93	5.250	17.98	0	0	3	4
16 Lincoln Continental	16.4	8	460.0	215	3.00	5.424	17.82	0	0	3	4
17 Chrysler Imperial	14.7	8	440.0	230	3.23	5.345	17.42	0	0	3	4

Data Frames - Builtin

R comes with a variety of built-in data sets that are very useful for getting used to data sets and how to manipulate them.

```
library(help="datasets")
```

```
# Gives detailed descriptions on available data sets
```

AirPassengers	Monthly Airline Passenger Numbers 1949-1960
BJsales	Sales Data with Leading Indicator
BOD	Biochemical Oxygen Demand
CO2	Carbon Dioxide Uptake in Grass Plants
ChickWeight	Weight versus age of chicks on different diets
DNase	Elisa assay of DNase
EuStockMarkets	Daily Closing Prices of Major European Stock Indices, 1991-1998
Formaldehyde	Determination of Formaldehyde
HairEyeColor	Hair and Eye Color of Statistics Students

```
help(mtcars) # Get details on a given data set
```

Data Frames - Builtin

R comes with a variety of built-in data sets that are very useful for getting used to data sets and how to manipulate them.

```
data(mtcars)
```

```
str(mtcars)
```

```
'data.frame':  32 obs. of  11 variables:
 $ mpg : num  21 21 22.8 21.4 18.7 18.1 14.3 24.4 22.8 19.2 ...
 $ cyl : num   6  6  4  6  8  6  8  4  4  6 ...
 $ disp: num  160 160 108 258 360 ...
 $ hp  : num  110 110  93 110 175 105 245  62  95 123 ...
 $ drat: num   3.9 3.9 3.85 3.08 3.15 2.76 3.21 3.69 3.92 3.92 ...
 $ wt  : num   2.62 2.88 2.32 3.21 3.44 ...
 $ qsec: num  16.5 17 18.6 19.4 17 ...
 $ vs  : num   0  0  1  1  0  1  0  1  1  1 ...
 $ am  : num   1  1  1  0  0  0  0  0  0  0 ...
 $ gear: num   4  4  4  3  3  3  3  4  4  4 ...
 $ carb: num   4  4  1  1  2  1  4  2  2  4 ...
```

```
nrow(mtcars) # How many rows does it have ?
```

```
[1] 32
```

```
ncol(mtcars) # How many columns are there ?
```

```
[1] 11
```

Data Frames - Accessing

There are various ways to select, remove, or exclude rows and columns

```
mtcars[,-11]
```

	mpg	cyl	disp	hp	drat	wt	qsec	vs	am	gear
Mazda RX4	21.0	6	160	110	3.90	2.620	16.46	0	1	4
Mazda RX4 Wag	21.0	6	160	110	3.90	2.875	17.02	0	1	4
Datsun 710	22.8	4	108	93	3.85	2.320	18.61	1	1	4

```
mtcars # Notice that carb is included
```

	mpg	cyl	disp	hp	drat	wt	qsec	vs	am	gear	carb
Mazda RX4	21.0	6	160.0	110	3.90	2.620	16.46	0	1	4	4
Mazda RX4 Wag	21.0	6	160.0	110	3.90	2.875	17.02	0	1	4	4
Datsun 710	22.8	4	108.0	93	3.85	2.320	18.61	1	1	4	1

Data Frames - Accessing

There are various ways to select, remove, or exclude rows and columns

```
mtcars[,-3:-5] # Print all columns except for columns 3 through 5
```

	mpg	cyl	wt	qsec	vs	am	gear	carb
Mazda RX4	21.0	6	2.620	16.46	0	1	4	0.6020600
Mazda RX4 Wag	21.0	6	2.875	17.02	0	1	4	0.6020600
Datsun 710	22.8	4	2.320	18.61	1	1	4	0.0000000

```
mtcars[,c(-3,-5)] # Print all columns except for columns 3 AND 5
```

	mpg	cyl	hp	wt	qsec	vs	am	gear	carb
Mazda RX4	21.0	6	110	2.620	16.46	0	1	4	0.6020600
Mazda RX4 Wag	21.0	6	110	2.875	17.02	0	1	4	0.6020600
Datsun 710	22.8	4	93	2.320	18.61	1	1	4	0.0000000

Data Frames - Accessing

There are various ways to select, remove, or exclude rows and columns

```
mtcars[mtcars$mpg >= 30.0,]
```

	mpg	cyl	disp	hp	drat	wt	qsec	vs	am	gear	carb
Fiat 128	32.4	4	78.7	66	4.08	2.200	19.47	1	1	4	1
Honda Civic	30.4	4	75.7	52	4.93	1.615	18.52	1	1	4	2
Toyota Corolla	33.9	4	71.1	65	4.22	1.835	19.90	1	1	4	1
Lotus Europa	30.4	4	95.1	113	3.77	1.513	16.90	1	1	5	2

```
mtcars[mtcars$mpg >= 30.0,2:6]
```

	mpg	cyl	disp	hp	drat
Fiat 128	32.4	4	78.7	66	4.08
Honda Civic	30.4	4	75.7	52	4.93
Toyota Corolla	33.9	4	71.1	65	4.22
Lotus Europa	30.4	4	95.1	113	3.77

```
mtcars[mtcars$mpg >= 30.0 & mtcars$cyl < 6,]
```

	mpg	cyl	disp	hp	drat	wt	qsec	vs	am	gear	carb
Fiat 128	32.4	4	78.7	66	4.08	2.200	19.47	1	1	4	1
Honda Civic	30.4	4	75.7	52	4.93	1.615	18.52	1	1	4	2
Toyota Corolla	33.9	4	71.1	65	4.22	1.835	19.90	1	1	4	1
Lotus Europa	30.4	4	95.1	113	3.77	1.513	16.90	1	1	5	2

Data Frames - Interrogating

Find all rows that correspond to Automatic and Count them

```
mtcars[mtcars$am==0,]  
      mpg  cyl  disp  hp drat   wt  qsec vs  am  gear  carb  
Hornet 4 Drive  21.4   6 258.0 110 3.08 3.215 19.44 1  0    3    1  
Hornet Sportabout 18.7   8 360.0 175 3.15 3.440 17.02 0  0    3    2  
Valiant        18.1   6 225.0 105 2.76 3.460 20.22 1  0    3    1  
Duster 360     14.3   8 360.0 245 3.21 3.570 15.84 0  0    3    4  
Merc 240D      24.4   4 146.7  62 3.69 3.190 20.00 1  0    4    2  
Merc 230       22.8   4 140.8  95 3.92 3.150 22.90 1  0    4    2  
..  
..
```

```
nrow(mtcars[mtcars$am == 0,])  
[1] 19
```

```
nrow(mtcars[mtcars$am == 1,])  
[1] 13
```

Data Frames - Interrogating

Extract all rows whose MPG value exceeds the mean MPG for the entire data frame

```
mtcars[mtcars$mpg > mean(mtcars$mpg),]
```

mpg	cyl	disp	hp	drat	wt	qsec	vs	am	gear	carb
Mazda RX4			21.0	6	160.0	110	3.90	2.620	16.46	0 1 4 4
Mazda RX4 Wag			21.0	6	160.0	110	3.90	2.875	17.02	0 1 4 4
Datsun 710			22.8	4	108.0	93	3.85	2.320	18.61	1 1 4 1
Hornet 4 Drive			21.4	6	258.0	110	3.08	3.215	19.44	1 0 3 1
Merc 240D			24.4	4	146.7	62	3.69	3.190	20.00	1 0 4 2
Merc 230			22.8	4	140.8	95	3.92	3.150	22.90	1 0 4 2
Fiat 128			32.4	4	78.7	66	4.08	2.200	19.47	1 1 4 1
Honda Civic			30.4	4	75.7	52	4.93	1.615	18.52	1 1 4 2
Toyota Corolla			33.9	4	71.1	65	4.22	1.835	19.90	1 1 4 1
Toyota Corona			21.5	4	120.1	97	3.70	2.465	20.01	1 0 3 1
Fiat X1-9			27.3	4	79.0	66	4.08	1.935	18.90	1 1 4 1
Porsche 914-2			26.0	4	120.3	91	4.43	2.140	16.70	0 1 5 2
Lotus Europa			30.4	4	95.1	113	3.77	1.513	16.90	1 1 5 2
Volvo 142E			21.4	4	121.0	109	4.11	2.780	18.60	1 1 4 2

Data Frames - Interrogating

Extract all rows whose MPG value exceeds the mean MPG for the entire data frame

```
# Find the quartiles for the MPG vector
```

```
quantile(mtcars$mpg)
      0%      25%      50%      75%     100%
10.400 15.425 19.200 22.800 33.900
```

```
# Now find the cars for which the MPG exceeds the 75% value:
```

```
mtcars[mtcars$mpg > quantile(mtcars$mpg)[4],]
```

	mpg	cyl	disp	hp	drat	wt	qsec	vs	am	gear	carb
Merc 240D	24.4	4	146.7	62	3.69	3.190	20.00	1	0	4	2
Fiat 128	32.4	4	78.7	66	4.08	2.200	19.47	1	1	4	1
Honda Civic	30.4	4	75.7	52	4.93	1.615	18.52	1	1	4	2
Toyota Corolla	33.9	4	71.1	65	4.22	1.835	19.90	1	1	4	1
Fiat X1-9	27.3	4	79.0	66	4.08	1.935	18.90	1	1	4	1
Porsche 914-2	26.0	4	120.3	91	4.43	2.140	16.70	0	1	5	2
Lotus Europa	30.4	4	95.1	113	3.77	1.513	16.90	1	1	5	2

Data Frames - Interrogating

What columns appear to be factors ? Variables with only a “few” different unique values perhaps ?

```
str(mtcars)
'data.frame':  32 obs. of  11 variables:
 $ mpg : num  21 21 22.8 21.4 18.7 18.1 14.3 24.4 22.8 19.2 ...
 $ cyl : num   6  6  4  6  8  6  8  4  4  6 ...
 $ disp: num  160 160 108 258 360 ...
 $ hp  : num  110 110 93 110 175 105 245 62 95 123 ...
 $ drat: num   3.9 3.9 3.85 3.08 3.15 2.76 3.21 3.69 3.92 3.92 ...
 $ wt  : num   2.62 2.88 2.32 3.21 3.44 ...
 $ qsec: num   16.5 17 18.6 19.4 17 ...
 $ vs  : num   0  0  1  1  0  1  0  1  1  1 ...
 $ am  : num   1  1  1  0  0  0  0  0  0  0 ...
 $ gear: num   4  4  4  3  3  3  3  4  4  4 ...
 $ carb: num   4  4  1  1  2  1  4  2  2  4 ...
```

```
unique(mtcars$am)  # Tells us what the unique values are
[1] 1 0
```

Data Frames - Factors

See how many unique values each column takes on

```
sapply(mtcars, function(x) length(unique(x)))
```

mpg	cyl	disp	hp	drat	wt	qsec	vs	am	gear	carb
25	3	27	22	22	29	30	2	2	3	6

If we summarize one of these potential factors right now, R will treat it as being purely numeric which we might not want

```
summary(mtcars$am)
```

Min.	1st Qu.	Median	Mean	3rd Qu.	Max.
0.0000	0.0000	0.0000	0.4062	1.0000	1.0000

So this really isn't helpful since we know that the "am" values are transmission types

```
mtcars$am <- factor(mtcars$am, levels = c(0,1), labels = c("Auto","Man") )
```

```
summary(mtcars$am)
```

Auto	Manu
19	13

Data Frames - Factors

We can add columns to a data frame. Let's say we want to create a new column called "mpgrate" that, based on the output of the quantile command, will have a rating of the that car's MPG in terms of "horrible", "bad", "good", "great".

The labels could be more scientific but this is still a good use case. There are a couple of ways to do this:

```
data(mtcars)    # Reload a "pure" copy of mtcars

mpgrate <- cut(mtcars$mpg,
              breaks = quantile(mtcars$mpg),
              labels=c("horrible", "Bad", "Good", "Great"), include.lowest=T)

mtcars <- cbind(mtcars,mpgrate)

-OR-

mtcars$mpgrate <- mpgrate    # The column just magically appears !
```


Data Frames - Factors


```
head(mtcars)
```

	mpg	cyl	disp	hp	drat	wt	qsec	vs	am	gear	carb	mpgrate
Mazda RX4	21.0	6	160	110	3.90	2.620	16.46	0	1	4	4	Good
Mazda RX4 Wag	21.0	6	160	110	3.90	2.875	17.02	0	1	4	4	Good
Datsun 710	22.8	4	108	93	3.85	2.320	18.61	1	1	4	1	Good
Hornet 4 Drive	21.4	6	258	110	3.08	3.215	19.44	1	0	3	1	Good
Hornet Sportabout	18.7	8	360	175	3.15	3.440	17.02	0	0	3	2	Bad
Valiant	18.1	6	225	105	2.76	3.460	20.22	1	0	3	1	Bad

```
library(lattice)
```

```
bwplot(~mpg|mpgrate,data=mtcars,layout=c(1,4))
```

Data Frames - Factors



Data Frames - transform()

You can also use the **transform()** command to change the types/classes of the columns

```
head(mtcars)
```

	mpg	cyl	disp	hp	drat	wt	qsec	vs	am	gear	carb
Mazda RX4	21.0	6	160	110	3.90	2.620	16.46	0	1	4	4
Mazda RX4 Wag	21.0	6	160	110	3.90	2.875	17.02	0	1	4	4
Datsun 710	22.8	4	108	93	3.85	2.320	18.61	1	1	4	1
Hornet 4 Drive	21.4	6	258	110	3.08	3.215	19.44	1	0	3	1
Hornet Sportabout	18.7	8	360	175	3.15	3.440	17.02	0	0	3	2
Valiant	18.1	6	225	105	2.76	3.460	20.22	1	0	3	1

```
transform(mtcars, wt = (wt*1000), qsec = round(qsec),  
          am = factor(am, labels=c("A", "M")))
```

	mpg	cyl	disp	hp	drat	wt	qsec	vs	am	gear	carb
Mazda RX4	21.0	6	160.0	110	3.90	2620	16	0	M	4	4
Mazda RX4 Wag	21.0	6	160.0	110	3.90	2875	17	0	M	4	4
Datsun 710	22.8	4	108.0	93	3.85	2320	19	1	M	4	1
Hornet 4 Drive	21.4	6	258.0	110	3.08	3215	19	1	A	3	1
Hornet Sportabout	18.7	8	360.0	175	3.15	3440	17	0	A	3	2

Data Frames - Reading CSV

Many times data will be read in from a comma delimited ,("CSV"), file exported from Excel. The file can be read from local storage or from the Web.

```
url <- "https://raw.githubusercontent.com/pittardsp/bios545r_spring_2018/master/SUPPORT/hsb2.csv"
```

```
data1 <- read.table(url,header=T,sep=",")
```

```
head(data1)
```

	gender	id	race	ses	schtyp	prgtype	read	write	math	science	socst
1	0	70	4	1	1	general	57	52	41	47	57
2	1	121	4	2	1	vocati	68	59	53	63	61
3	0	86	4	3	1	general	44	33	54	58	31
4	0	141	4	3	1	vocati	63	44	47	53	56
5	0	172	4	2	1	academic	47	52	57	53	61
6	0	113	4	2	1	academic	44	52	51	63	61

Motivations

dplyr is an add on package designed to efficiently transform and summarize tabular data such as data frames. The package has a number of functions ("verbs") that perform a number of data manipulation tasks:

- Filtering rows
- Select specific columns
- Re-ordering or arranging rows
- Summarizing and aggregating data

One of the unique strengths of **dplyr** is that it implements what is known as a Split-Apply-Combine technique that we will explore in this session.

The `dyplr` function can also be used with the **magrittr** package for setting up workflows or pipelines to process data.

Motivations

- **dplyr** is designed to work with data frames but it can also connect to relational databases that are locally or remotely available.
- Access to data frames or databases is accomplished using the same set of tools. You don't have to use different commands.
- Relative to databases you use the “verbs” provided with dplyr that in turn are translated into the appropriate SQL statements necessary to interact with the databases.

How to Install dplyr ?

```
# install the package
```

```
install.packages("dplyr")
```

```
install.packages("readr")    # Get's the equivalent to data.table's fread p
```

```
# Loads the package
```

```
library(dplyr)
```

```
# Launches a browser to explore
```

```
browseVignettes(package = "dplyr")
```

Motivations

This slide deck references “Becoming a data ninja with dplyr” <https://speakerdeck.com/dpastoor/becoming-a-data-ninja-with-dplyr> and the dplyr tutorial http://genomicsclass.github.io/book/pages/dplyr_tutorial.html

Motivations

- A data frame is a set of columns. Every column is same length but of possibly different types.
- It has characteristics of both a matrix, (each row is the same data type),
- Each column can be a different data type
- Bracket notation offers a convenient way to search through the data frame

Motivations

```
head(mtcars, 12)
```

	mpg	cyl	disp	hp	drat	wt	qsec	vs	am	gear	carb
Mazda RX4	21.0	6	160.0	110	3.90	2.620	16.46	0	1	4	4
Mazda RX4 Wag	21.0	6	160.0	110	3.90	2.875	17.02	0	1	4	4
Datsun 710	22.8	4	108.0	93	3.85	2.320	18.61	1	1	4	1
Hornet 4 Drive	21.4	6	258.0	110	3.08	3.215	19.44	1	0	3	1
Hornet Sportabout	18.7	8	360.0	175	3.15	3.440	17.02	0	0	3	2
Valiant	18.1	6	225.0	105	2.76	3.460	20.22	1	0	3	1
Duster 360	14.3	8	360.0	245	3.21	3.570	15.84	0	0	3	4
Merc 240D	24.4	4	146.7	62	3.69	3.190	20.00	1	0	4	2
Merc 230	22.8	4	140.8	95	3.92	3.150	22.90	1	0	4	2
Merc 280	19.2	6	167.6	123	3.92	3.440	18.30	1	0	4	4
Merc 280C	17.8	6	167.6	123	3.92	3.440	18.90	1	0	4	4
Merc 450SE	16.4	8	275.8	180	3.07	4.070	17.40	0	0	3	3

Motivations

There are some common activities associated with a data frame:

- filter - find observations satisfying some condition(s)
- select - selecting specific columns by name
- mutate - adding new columns or changing existing ones
- arrange - reorder or sort the rows
- summarize - do some aggregation or summary by groups

Motivations

```
df <- data.frame(id = 1:5,  
                 gender = c("MALE", "MALE", "FEMALE", "MALE", "FEMALE"),  
                 age = c(70, 76, 60, 64, 68))
```

ID	GENDER	AGE
1	MALE	70
2	MALE	76
3	FEMALE	60
4	MALE	64
5	FEMALE	68

Filter

```
filter(df, gender == "FEMALE")
```

```
  id gender age  
1  3 FEMALE  60  
2  5 FEMALE  68
```

ID	GENDER	AGE
1	MALE	70
2	MALE	76
3	FEMALE	60
4	MALE	64
5	FEMALE	68

ID	GENDER	AGE
3	FEMALE	60
5	FEMALE	68

Filter

```
filter(df, id %in% c(1,3,5))
```

```
  id gender age  
1  1  MALE  70  
2  3 FEMALE  60  
3  5 FEMALE  68
```

ID	GENDER	AGE
1	MALE	70
2	MALE	76
3	FEMALE	60
4	MALE	64
5	FEMALE	68

ID	GENDER	AGE
1	MALE	70
3	FEMALE	60
5	FEMALE	68

Mutate

Mutate is used to add or remove columns in a data frame

```
mutate(df, meanage = mean(age))
```

```
  id gender age meanage
1  1  MALE  70    67.6
2  2  MALE  76    67.6
3  3 FEMALE  60    67.6
4  4  MALE  64    67.6
5  5 FEMALE  68    67.6
```

ID	GENDER	AGE
1	MALE	70
2	MALE	76
3	FEMALE	60
4	MALE	64
5	FEMALE	68

ID	GENDER	AGE	MEANWT
1	MALE	70	67.6
2	MALE	76	67.6
3	FEMALE	60	67.6
4	MALE	64	67.6
5	FEMALE	68	67.6

Mutate

Here we create a new column designed to tell us if a given observation has an age that is greater than or equal to the average age.


We create a variable called `old_young` and assign a value of “Y” if they are above the mean age and a value of “N” if they are not.

```
mutate(df, old_young = ifelse(df$age >= mean(df$age), "Y", "N"))
```

	id	gender	age	old_young
1	1	MALE	70	Y
2	2	MALE	76	Y
3	3	FEMALE	60	N
4	4	MALE	64	N
5	5	FEMALE	68	Y

Mutate

```
tmp <- mutate(df, color = ifelse(age > mean(age), "red", "blue"))
plot(tmp$age, col=tmp$color, type="p", pch=19, main="Ages", ylab="Age")
grid()
abline(h=mean(tmp$age), lty=2)
legend("topright", c("Above Avg", "Below Avg"), col=c("red", "blue"), pch=19)
```



Arrange

Use arrange for sorting the data frame by a column(s)

```
# Sort df by age from highest to lowest
```

```
arrange(df, desc(age))
```

	id	gender	age
1	2	MALE	76
2	1	MALE	70
3	5	FEMALE	68
4	4	MALE	64
5	3	FEMALE	60

```
# Sort df by gender (alphabetically) and then by age  
# from highest to lowest
```

```
arrange(df, gender, desc(age))
```

	id	gender	age
1	5	FEMALE	68
2	3	FEMALE	60
3	2	MALE	76
4	1	MALE	70
5	4	MALE	64

Select

Select allows us to select groups of columns from a data frame

```
select(df,gender,id,age) # Reorder the columns
```

```
  gender id age
1  MALE  1  70
2  MALE  2  76
3 FEMALE  3  60
4  MALE  4  64
5 FEMALE  5  68
```

```
select(df,-age) # Select all but the age column
```

```
  id gender
1  1  MALE
2  2  MALE
3  3 FEMALE
4  4  MALE
5  5 FEMALE
```

```
select(df,id:age) # Can use : to select a range
```

```
  id gender age
1  1  MALE  70
2  2  MALE  76
3  3 FEMALE  60
4  4  MALE  64
5  5 FEMALE  68
```

Select

You can select by regular expressions or numeric patterns

```
library(ggplot2)
data(diamonds)
names(diamonds)
[1] "carat" "cut" "color" "clarity" "depth" "table" "price"
[8] "x" "y" "z"
```

```
head(select(diamonds,starts_with("c")))
```

	carat	cut	color	clarity
1	0.23	Ideal	E	SI2
2	0.21	Premium	E	SI1
3	0.23	Good	E	VS1
4	0.29	Premium	I	VS2
5	0.31	Good	J	SI2
6	0.24	Very Good	J	VVS2

```
head(select(diamonds,ends_with("t")))
```

	carat	cut
1	0.23	Ideal
2	0.21	Premium
3	0.23	Good
4	0.29	Premium
5	0.31	Good
6	0.24	Very Good

Select

You can select by regular expressions or numeric patterns

```
testdf <- expand.grid(m_1=seq(60,70,10),age=c(25,32),m_2=seq(50,60,10))
```

```
head(testdf, 4)
```

	m_1	age	m_2
1	60	25	50
2	70	25	50
3	60	32	50
4	70	32	50

```
head( select(testdf,matches("_")) ,2)
```

	m_1	m_2
1	60	50
2	70	50

```
head( select(testdf,contains("_"), 2)
```

	m_1	m_2
1	60	50
2	70	50

```
head( select(testdf,num_range("m_",1:2)), 2)
```

	m_1	m_2
1	60	50
2	70	50

group_by

group_by let's you organize a data frame by some factor or grouping variable

```
df
  id gender age
1  1   MALE  70
2  2   MALE  76
3  3 FEMALE  60
4  4   MALE  64
5  5 FEMALE  68
```

```
group_by(df,gender)    # Hmm. Did this really do anything ?
```

```
Source: local data frame [5 x 3]
Groups: gender
```

```
  id gender age
1  1   MALE  70
2  2   MALE  76
3  3 FEMALE  60
4  4   MALE  64
5  5 FEMALE  68
```

group_by

group_by let's you organize a data frame by some factor or grouping variable

```
df
```

```
  id gender age
1  1  MALE  70
2  2  MALE  76
3  3 FEMALE  60
4  4  MALE  64
5  5 FEMALE  68
```

```
( gdf <- group_by(df,gender)    # Hmm. Did this really do anything ?
```

```
Source: local data frame [5 x 3]
```

```
Groups: gender
```

```
  id gender age
1  1  MALE  70
2  2  MALE  76
3  3 FEMALE  60
4  4  MALE  64
5  5 FEMALE  68
```

Summarize

```
summarize(group_by(df,gender),total=n())
```

Source: local data frame [2 x 2]

```
gender total
1 FEMALE    2
2  MALE     3
```

ID	GENDER	AGE
1	MALE	70
2	MALE	76
3	FEMALE	60
4	MALE	64
5	FEMALE	68

GENDER	TOTAL
FEMALE	2
MALE	3

Summarize

```
summarize(group_by(df,gender),av_age=mean(age))
```

Source: local data frame [2 x 2]

```
gender av_age
1 FEMALE    64
2  MALE     70
```

ID	GENDER	AGE
1	MALE	70
2	MALE	76
3	FEMALE	60
4	MALE	64
5	FEMALE	68

GENDER	AV_AGE
FEMALE	64
MALE	70

Summarize

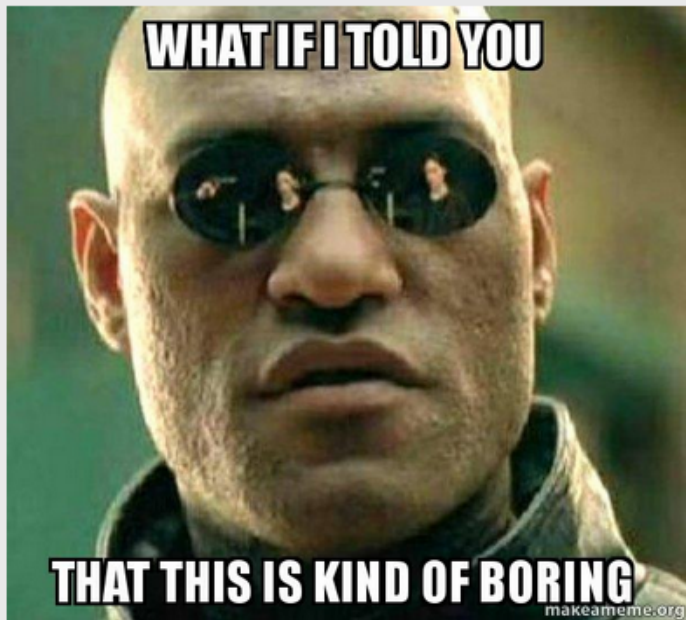
```
summarize(group_by(df,gender),av_age=mean(age),total=n())
```

Source: local data frame [2 x 3]

```
gender av_age total
1 FEMALE      64     2
2  MALE      70     3
```

ID	GENDER	AGE
1	MALE	70
2	MALE	76
3	FEMALE	60
4	MALE	64
5	FEMALE	68

GENDER	AV_AGE	TOTAL
FEMALE	64	2
MALE	70	3



Split -> Apply -> Combine

Split -> Apply -> Combine

group_by

ID	GENDER	AGE
1	MALE	70
2	MALE	76
3	FEMALE	60
4	MALE	64
5	FEMALE	68

ID	GENDER	AGE
1	MALE	70
2	MALE	76
4	MALE	64

ID	GENDER	AGE
3	FEMALE	60
5	FEMALE	68

AVG
70

AVG
64

ID	GENDER	AVG
1	MALE	70
2	FEMALE	64

Split -> Apply -> Combine

But do you really need dplyr to do this ? No but it makes it a lot easier

df

```
  id gender age
1  1  MALE  70
2  2  MALE  76
3  3 FEMALE  60
4  4  MALE  64
5  5 FEMALE  68
```

```
tapply(df$age,df$gender,mean)  # tapply function
```

```
FEMALE  MALE
    64    70
```

```
aggregate(age~gender,data=df,mean) # aggregate works also
```

```
gender age
1 FEMALE  64
2  MALE   70
```

```
lapply(split(df,df$gender),function(x) mean(x$age)) # complicated
```

```
$FEMALE
```

```
[1] 64
```

```
$MALE
```

```
[1] 70
```

Split -> Apply -> Combine: Chaining

- Before moving forward let us consider the “pipe” operator that is included with the **magrittr** package. This is used to make it possible to “pipe” the results of one command into another command and so on.
- The inspiration for this comes from the UNIX/LINUX operating system where pipes are used all the time. So in effect using “pipes” is nothing new in the world of research computation.
- **Warning:** Once you get used to pipes it is hard to go back to not using them

Split -> Apply -> Combine: Chaining

When you load the dplyr package it in turn loads the necessary packages for supporting the piping capability. Let's use the mtcars data frame to illustrate the basics of the piping mechanism as used by dplyr.

Here we will select the mpg and am column from mtcars and view the top 5 rows.

```
head(select(mtcars, mpg, am))
```

	mpg	am
Mazda RX4	21.0	1
Mazda RX4 Wag	21.0	1
Datsun 710	22.8	1
Hornet 4 Drive	21.4	0
Hornet Sportabout	18.7	0
Valiant	18.1	0

Split -> Apply -> Combine: Chaining

Here we will select the mpg and am column from mtcars and view the top 5 rows but using dplyr and the piping operator. Instead of nesting functions (reading from the inside to the outside), the idea of of piping is to read the functions from left to right.

```
mtcars %>% select(mpg, am) %>% head
```

	mpg	am
Mazda RX4	21.0	1
Mazda RX4 Wag	21.0	1
Datsun 710	22.8	1
Hornet 4 Drive	21.4	0
Hornet Sportabout	18.7	0
Valiant	18.1	0

Split -> Apply -> Combine: Chaining

What about this ? We can chain together the output of one command to the input of another !

```
df %>% group_by(gender) %>% summarize(avg=mean(age))
```

Source: local data frame [2 x 2]

	gender	avg
1	FEMALE	64
2	MALE	70

```
df %>% group_by(gender) %>% summarize(avg=mean(age),total=n())
```

Source: local data frame [2 x 3]

	gender	avg	total
1	FEMALE	64	2
2	MALE	70	3

```
df %>% filter(gender == "MALE") %>% summarize(med_age=median(age))
```

	med_age
1	70

Split -> Apply -> Combine: Chaining

What about this ? We can chain together the output of one command to the input of another !

```
df %>% filter(gender == "MALE") %>% summarize(med_age=median(age))
```

	med_age
1	70

df

ID	GENDER	AGE
1	MALE	70
2	MALE	76
3	FEMALE	60
4	MALE	64
5	FEMALE	68

filter

ID	GENDER	AGE
1	MALE	70
2	MALE	76
4	MALE	64

summarize

<u>med_age</u>
70

Split -> Apply -> Combine: Chaining

Using the built in mtcars dataframe filter out records where the wt is greater than 3.3 tons.

Then create a column called ab_be (Y or N) that indicates whether that observation's mpg is greater (or not) than the average mpg for the filtered set.

Then present the average mpg for each group

```
mtcars %>% filter(wt > 3.3) %>%  
  mutate(ab_be=ifelse(mpg > mean(mpg),"Y","N")) %>%  
  group_by(ab_be) %>% summarize(mean_mpg=mean(mpg))
```

Source: local data frame [2 x 2]

	ab_be	mean_mpg
1	N	13.77778
2	Y	18.10000

Split -> Apply -> Combine: Chaining

Using the built in mtcars dataframe filter out records where the wt is greater than 3.3 tons.

```
mtcars %>% filter(wt > 3.3)
```

	mpg	cyl	disp	hp	drat	wt	qsec	vs	am	gear	carb
1	18.7	8	360.0	175	3.15	3.440	17.02	0	0	3	2
2	18.1	6	225.0	105	2.76	3.460	20.22	1	0	3	1
3	14.3	8	360.0	245	3.21	3.570	15.84	0	0	3	4
4	19.2	6	167.6	123	3.92	3.440	18.30	1	0	4	4
5	17.8	6	167.6	123	3.92	3.440	18.90	1	0	4	4
6	16.4	8	275.8	180	3.07	4.070	17.40	0	0	3	3
7	17.3	8	275.8	180	3.07	3.730	17.60	0	0	3	3
8	15.2	8	275.8	180	3.07	3.780	18.00	0	0	3	3
9	10.4	8	472.0	205	2.93	5.250	17.98	0	0	3	4
10	10.4	8	460.0	215	3.00	5.424	17.82	0	0	3	4
11	14.7	8	440.0	230	3.23	5.345	17.42	0	0	3	4
12	15.5	8	318.0	150	2.76	3.520	16.87	0	0	3	2
13	15.2	8	304.0	150	3.15	3.435	17.30	0	0	3	2
14	13.3	8	350.0	245	3.73	3.840	15.41	0	0	3	4
15	19.2	8	400.0	175	3.08	3.845	17.05	0	0	3	2
16	15.0	8	301.0	335	3.54	3.570	14.60	0	1	5	8

Split -> Apply -> Combine: Chaining

Create a column called `ab_be` (Y or N) that indicates whether that observation's mpg is greater (or not) than the average mpg for the filtered set.

```
mtcars %>% filter(wt > 3.3) %>%  
  mutate(ab_be=ifelse(mpg > mean(mpg),"Y","N"))  
  mpg cyl  disp  hp drat   wt  qsec vs am gear carb ab_be  
1  18.7   8 360.0 175 3.15 3.440 17.02  0  0   3   2    Y  
2  18.1   6 225.0 105 2.76 3.460 20.22  1  0   3   1    Y  
3  14.3   8 360.0 245 3.21 3.570 15.84  0  0   3   4    N  
4  19.2   6 167.6 123 3.92 3.440 18.30  1  0   4   4    Y  
5  17.8   6 167.6 123 3.92 3.440 18.90  1  0   4   4    Y  
6  16.4   8 275.8 180 3.07 4.070 17.40  0  0   3   3    Y  
7  17.3   8 275.8 180 3.07 3.730 17.60  0  0   3   3    Y  
8  15.2   8 275.8 180 3.07 3.780 18.00  0  0   3   3    N  
9  10.4   8 472.0 205 2.93 5.250 17.98  0  0   3   4    N  
10 10.4   8 460.0 215 3.00 5.424 17.82  0  0   3   4    N  
11 14.7   8 440.0 230 3.23 5.345 17.42  0  0   3   4    N  
12 15.5   8 318.0 150 2.76 3.520 16.87  0  0   3   2    N  
13 15.2   8 304.0 150 3.15 3.435 17.30  0  0   3   2    N  
14 13.3   8 350.0 245 3.73 3.840 15.41  0  0   3   4    N  
15 19.2   8 400.0 175 3.08 3.845 17.05  0  0   3   2    Y  
16 15.0   8 301.0 335 3.54 3.570 14.60  0  1   5   8    N
```

Split -> Apply -> Combine: Chaining

Then present the average mpg for each group as defined by ab_be

```
mtcars %>% filter(wt > 3.3) %>%  
  mutate(ab_be=ifelse(mpg > mean(mpg),"Y","N")) %>%  
  group_by(ab_be) %>% summarize(mean_mpg=mean(mpg))
```


Source: local data frame [2 x 2]

	ab_be	mean_mpg
1	N	13.77778
2	Y	18.10000

Split -> Apply -> Combine: Chaining

This could then be chained to the ggplot command

```
mtcars %>% filter(wt > 3.3) %>%  
  mutate(ab_be=ifelse(mpg > mean(mpg),"Y","N")) %>%  
  group_by(ab_be) %>% summarize(mean_mpg=mean(mpg)) %>%  
  ggplot(aes(x=ab_be,y=mean_mpg)) + geom_bar(stat="identity") +  
  ggtitle("Mean MPG") + labs(x = "ab_be", y = "Mean MPG")
```



Large Files

Let's read in the file combined_wiki.txt.gz

```
library(readr)
```

```
dt <- read_delim("combined_wiki.zip",delim=" ")
```

```
nrow(dt)
```

```
[1] 31164567
```

```
head(dt,5)
```

	proj	page	acc	bytes
1:	aa.b	Main_Page	1	5565
2:	aa.b	MediaWiki:Image_sample	1	5179
3:	aa.b	MediaWiki:Upload_source_file	1	5195
4:	aa.b	Wikibooks:Privacy_policy	1	4925
5:	aa.d	MediaWiki:Group-abusefilter-member	1	4912

Large Files

Using dplyr commands, summarize the mean number of bytes (in megabytes) per unique project page and sort the resulting table in descending order by the average in megabytes.

```
nrow(dt)
[1] 31164567
```

```
head(dt,5)
```

	proj	page	acc	bytes
1:	aa.b	Main_Page	1	5565
2:	aa.b	MediaWiki:Image_sample	1	5179
3:	aa.b	MediaWiki:Upload_source_file	1	5195
4:	aa.b	Wikibooks:Privacy_policy	1	4925
5:	aa.d	MediaWiki:Group-abusefilter-member	1	4912

Split -> Apply -> Combine: Chaining

Using dplyr commands, summarize the mean number of bytes (in megabytes) per unique project page and sort the resulting table in descending order by the average in megabytes.

```
dt %>% mutate(MB=bytes/1000000) %>%  
  group_by(proj)%>%  
  summarize(avg=round(mean(MB),2)) %>%  
  arrange(desc(avg))
```

Source: local data table [1,266 x 2] # Note we have 1,266 rows

	V1	avg
1	en.mw	77518.22
2	ja.mw	9126.98
3	fr.mw	2020.45
4	ru.mw	1311.16
5	de.mw	1214.59
6	es.mw	1187.93
7	it.mw	472.27
8	zh.mw	374.91
9	ko.mw	234.63
10	pt.mw	207.78

Split -> Apply -> Combine: Chaining

Using dplyr commands, summarize the mean number of bytes (in megabytes) per unique project page and sort the resulting table in descending order by the average in megabytes.

```
system.time( dt %>% mutate(MB=bytes/1000000) %>%  
              group_by(proj)%>%  
              summarize(avg=round(mean(MB),2)) %>%  
              arrange(desc(avg)) )
```

user	system	elapsed
1.93	2.58	7.79

dply vs Native R Commands

How long with this take using the standard native R commands ? First we create a function so we can easily time things. This example also assumes that we have read in the data using the native R command **read.csv** file command. That is we are not benefitting from the conveniences offered by the `data.table` command.

```
myaggre <- function(df) {  
  df$bytes <- round(df$bytes/1000000,2)  
  hold <- aggregate(bytes~proj,df,mean)  
  hold <- hold[order(-hold$bytes),]  
  return(hold)  
}
```

```
system.time( myaggre(df))  
   user  system elapsed  
351.826  11.120  378.115
```

dplyr additional commands

Other activities are possible

```
mtcars %>% sample_n(2) # Sample 2 records from a data frame
```

	mpg	cyl	disp	hp	drat	wt	qsec	vs	am	gear	carb
Mazda RX4 Wag	21.0	6	160.0	110	3.90	2.875	17.02	0	1	4	4
Merc 280	19.2	6	167.6	123	3.92	3.440	18.30	1	0	4	4

```
# Sample 2 records from each cylinder group
```

```
mtcars %>% group_by(cyl) %>% do(sample_n(.,2))
```

```
Source: local data frame [6 x 11]
```

```
Groups: cyl
```

	mpg	cyl	disp	hp	drat	wt	qsec	vs	am	gear	carb
1	21.4	4	121.0	109	4.11	2.780	18.60	1	1	4	2
2	27.3	4	79.0	66	4.08	1.935	18.90	1	1	4	1
3	19.7	6	145.0	175	3.62	2.770	15.50	0	1	5	6
4	18.1	6	225.0	105	2.76	3.460	20.22	1	0	3	1
5	17.3	8	275.8	180	3.07	3.730	17.60	0	0	3	3
6	19.2	8	400.0	175	3.08	3.845	17.05	0	0	3	2

dplyr additional commands

Other activities are possible. You can use “do” to perform arbitrary computation, returning either a data frame or arbitrary objects which will be stored in a list. This is particularly useful when working with models

```
by_cyl <- group_by(mtcars, cyl)
models <- by_cyl %>% do(mod = lm(mpg ~ disp, data = .))
```

Source: local data frame [3 x 2]

Groups: <by row>

	cyl	mod
1	4	<S3:lm>
2	6	<S3:lm>
3	8	<S3:lm>

```
summarise(models, rsq = summary(mod)$r.squared)
```

Source: local data frame [3 x 1]

	rsq
1	0.64840514
2	0.01062604
3	0.27015777

Here is a one liner that does the above

Joining data frames

```
idatime <- data.frame(id=rep(1:3,each=2),time=rep(0:1,each=3))
```

	id	time
1	1	0
2	1	0
3	2	0
4	2	1
5	3	1
6	3	1

```
idawt <- data.frame(id=c(1,2,4),wt=c(110,130,115))
```

	id	wt
1	1	110
2	2	130
3	4	115

Inner joins - inner_join(x,y)

Will return all rows from x where there are matching values in y, and all columns from x and y

idatime		idawt		inner_join(idatime, idawt)			inner_join(idawt, idatime)		
id	time	id	wt	id	time	wt	id	wt	time
1	0	1	110	1	0	110	1	110	0
1	0	2	130	1	0	110	1	110	0
2	0	4	115	2	0	130	2	130	0
2	1			2	1	130	2	130	1
3	1								
3	1								

Inner joins - inner_join(x,y)

Will return all rows from x where there are matching values in y, and all columns from x and y

```
inner_join(idatime,idawt)
```

Joining by: "id"

	id	time	wt
1	1	0	110
2	1	0	110
3	2	0	130
4	2	1	130

Joining data frames - left_join(x,y)

return all rows from x, and all columns from x and y

idatime

id	time
1	0
1	0
2	0
2	1
3	1
3	1

idawt

id	wt
1	110
2	130
4	115

left_join(idatime, idawt)

id	time	wt
1	0	110
1	0	110
2	0	130
2	1	130
3	1	NA
3	1	NA

left_join(idawt, idatime)

id	wt	time
1	110	0
1	110	0
2	130	0
2	130	1
4	115	NA

Joining data frames - anti_join(x,y)

returns all rows from x where there are not any matching values in y, keeping just the columns from x

idatime		idawt		anti_join(idatime, idawt)		anti_join(idawt, idatime)	
id	time	id	wt	id	time	id	wt
1	0	1	110	3	1	4	115
1	0	2	130	3	1		
2	0	4	115				
2	1						
3	1						
3	1						

Joining data frames

idatime	
id	time
1	0
1	0
2	0
2	1
3	1
3	1

idawt	
id	wt
1	110
2	130
4	115

`semi_join(idatime, idawt)`

id	time
1	0
1	0
2	0
2	1

`semi_join(idawt, idatime)`

id	wt
1	110
2	130

Joining data frames

Return all rows from x where there are matching values in y, keeping just columns from x

idatime		idawt		semi_join(idatime, idawt)		semi_join(idawt, idatime)	
id	time	id	wt	id	time	id	wt
1	0	1	110	1	0	1	110
1	0	2	130	1	0	2	130
2	0	4	115	2	0		
2	1			2	1		
3	1						
3	1						

Exercises

Now it is your turn. Let's do some exercises

```
url <- "https://raw.githubusercontent.com/pittardsp/bios545r_spring_2018/master/SUPPORT/msleep_ggplot2.csv"
```

```
library(readr)
```

```
download.file(url, "msleep_ggplot2.csv")
```

```
msleep <- read_csv("msleep_ggplot2.csv")
```

```
names(msleep)
```

[1]	"name"	"genus"	"vore"	"order"	"conservation"
[6]	"sleep_total"	"sleep_rem"	"sleep_cycle"	"awake"	"brainwt"
[11]	"bodywt"				

Exercises

Here is a description of the columns / variables

COLUMN NAME	DESCRIPTION
-------------	-------------

name	common name
genus	taxonomic rank
vore	carnivore, omnivore or herbivore?
order	taxonomic rank
conservation	the conservation status of the mammal
sleep_total	total amount of sleep, in hours
sleep_rem	rem sleep, in hours
sleep_cycle	length of sleep cycle, in hours
awake	amount of time spent awake, in hours
brainwt	brain weight in kilograms
bodywt	body weight in kilograms

Exercises

Using the `msleep` data frame let's do the following activities and answer some questions. Try to use the chaining operator:

- Select the `name` and `sleep_total` columns
- Using the colon operator select all columns between `name` and `order`
- Select all columns that begin with "sl"
- Filter the data frame to find only rows with a `sleep_total` ≥ 16
- Filter the rows for mammals that sleep a total of more than 16 hours and have a body weight of greater than 1 kilogram
- Arrange the data frame using the **order** column

Exercises

For these you will need to use the chaining operator:

- Select three columns from `msleep`, arrange the rows by the taxonomic order and then arrange the rows by `sleep_total`. Finally show the head of the final data frame
- Same as above, except here we filter the rows for mammals that sleep for 16 or more hours instead of showing the head of the final data frame
- Use the **`mutate`** function to create a new column called `rem_proportion` which is the ratio of rem sleep to total amount of sleep

Exercises

- Use the summarise() function: to compute the average number of hours of sleep, apply the mean() function to the column sleep_total and call the summary value avg_sleep.
- Group the msleep data frame by taxonomic order and then summarize the mean sleep total
- Same as above except summarize the mean, max, and min sleep total